

Table Of Contents

- What's New
- Installation
- Contributing to pandas
- Package overview
- 10 Minutes to pandas
 - Object Creation
 - Viewing Data
 - Selection
 - Getting
 - Selection by Label
 - Selection by Position
 - Boolean Indexing
 - Setting
 - Missing Data
 - Operations
 - Stats
 - Apply
 - Histogramming
 - String Methods
 - Merge
 - Concat
 - Join
 - Append
 - Grouping
 - Reshaping
 - Stack
 - Pivot Tables
 - Time Series
 - Categoricals
 - Plotting
 - Getting Data In/Out
 - CSV
 - HDF5
 - Excel
 - Gotchas
- Tutorials
- Cookbook
- Intro to Data Structures
- Essential Basic Functionality
- Working with Text Data
- Options and Settings
- Indexing and Selecting Data
- MultIndex / Advanced Indexing
- Computational tools
- Working with missing data
- Group By: split-apply-combine
- Merge, join, and concatenate
- Reshaping and Pivot Tables
- Time Series / Date functionality
- Time Deltas
- Categorical Data
- Visualization
- Styling
- IO Tools (Text, CSV, HDF5, ...)
- Enhancing Performance
- Sparse data structures
- Frequently Asked Questions (FAQ)
- rpy2 / R interface
- pandas Ecosystem
- Comparison with R / R libraries
- Comparison with SQL
- Comparison with SAS
- Comparison with Stata
- API Reference
- Developer
- Internals
- Extending Pandas
- Release Notes

10 Minutes to pandas

This is a short introduction to pandas, geared mainly for new users. You can see more complex recipes in the [Cookbook](#).

Customarily, we import as follows:

```
In [1]: import pandas as pd
In [2]: import numpy as np
In [3]: import matplotlib.pyplot as plt
```

Object Creation

See the [Data Structure Intro section](#).

Creating a [Series](#) by passing a list of values, letting pandas create a default integer index:

```
In [4]: s = pd.Series([1,3,5,np.nan,6,8])
In [5]: s
Out[5]:
0    1.0
1    3.0
2    5.0
3    NaN
4    6.0
5    8.0
dtype: float64
```

Creating a [DataFrame](#) by passing a NumPy array, with a datetime index and labeled columns:

```
In [6]: dates = pd.date_range('20130101', periods=6)
In [7]: dates
Out[7]:
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')
In [8]: df = pd.DataFrame(np.random.randn(6,4), index=dates, columns=list('ABCD'))
In [9]: df
Out[9]:
   A      B      C      D
2013-01-01  0.469112 -0.282863 -1.509059 -1.135632
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804
2013-01-04  0.721555 -0.706771 -1.039575  0.271860
2013-01-05 -0.424972  0.567020  0.276232 -1.087401
2013-01-06 -0.673690  0.113648 -1.478427  0.524988
```

Creating a DataFrame by passing a dict of objects that can be converted to series-like.

```
In [10]: df2 = pd.DataFrame({ 'A' : 1.,
                           ....: 'B' : pd.Timestamp('20130102'),
                           ....: 'C' : pd.Series(1,index=list(range(4)),dtype='float32'),
                           ....: 'D' : np.array([3] * 4,dtype='int32'),
                           ....: 'E' : pd.Categorical(["test","train","test","train"]),
                           ....: 'F' : 'foo' })
In [11]: df2      Add another column G containing your roll
Out[11]: number
   A      B      C      D      E      F
0  1.0  2013-01-02  1.0  3  test  foo
1  1.0  2013-01-02  1.0  3  train  foo
2  1.0  2013-01-02  1.0  3  test  foo
3  1.0  2013-01-02  1.0  3  train  foo
```

The columns of the resulting DataFrame have different [dtypes](#).

Go

Enter search terms or a module, class or function name.

In [12]: df2.dtypes**Out[12]:**

A float64
B datetime64[ns]
C float32
D int32
E category
F object
dtype: object

If you're using IPython, tab completion for column names (as well as public attributes) is automatically enabled. Here's a subset of the attributes that will be completed:

In [13]: df2.<TAB>

df2.A df2.bool
df2.abs df2.boxplot
df2.add df2.C
df2.add_prefix df2.clip
df2.add_suffix df2.clip_lower
df2.align df2.clip_upper
df2.all df2.columns
df2.any df2.combine
df2.append df2.combine_first
df2.apply df2.compound
df2.applomap df2.consolidate
df2.D

As you can see, the columns A, B, C, and D are automatically tab completed. E is there as well; the rest of the attributes have been truncated for brevity.

Viewing Data

See the [Basics](#) section.

Here is how to view the top and bottom rows of the frame:

In [14]: df.head() Print first 4 rows only**Out[14]:**

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

In [15]: df.tail(3) Print last 5 rows only**Out[15]:**

	A	B	C	D
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

Display the index, columns, and the underlying NumPy data:

In [16]: df.index**Out[16]:**

DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
'2013-01-05', '2013-01-06'],
dtype='datetime64[ns]', freq='D')

In [17]: df.columns

Out[17]: Index(['A', 'B', 'C', 'D'], dtype='object')

In [18]: df.values**Out[18]:**

```
array([[ 0.4691, -0.2829, -1.5091, -1.1356],  
       [ 1.2121, -0.1732,  0.1192, -1.0442],  
       [-0.8618, -2.1046, -0.4949,  1.0718],  
       [ 0.7216, -0.7068, -1.0396,  0.2719],  
       [-0.425 ,  0.567 ,  0.2762, -1.0874],  
       [-0.6737,  0.1136, -1.4784,  0.525 ]])
```

[describe\(\)](#) shows a quick statistic summary of your data:

```
In [19]: df.describe()
Out[19]:
      A      B      C      D
count 6.000000 6.000000 6.000000 6.000000
mean 0.073711 -0.431125 -0.687758 -0.233103
std 0.843157 0.922818 0.779887 0.973118
min -0.861849 -2.104569 -1.509059 -1.135632
25% -0.611510 -0.600794 -1.368714 -1.076610
50% 0.022070 -0.228039 -0.767252 -0.386188
75% 0.658444 0.041933 -0.034326 0.461706
max 1.212112 0.567020 0.276232 1.071804
```

Transposing your data:

```
In [20]: df.T
Out[20]:
2013-01-01 2013-01-02 2013-01-03 2013-01-04 2013-01-05 2013-01-06
A 0.469112 1.212112 -0.861849 0.721555 -0.424972 -0.673690
B -0.282863 -0.173215 -2.104569 -0.706771 0.567020 0.113648
C -1.509059 0.119209 -0.494929 -1.039575 0.276232 -1.478427
D -1.135632 -1.044236 1.071804 0.271860 -1.087401 0.524988
```

Sorting by an axis:

```
In [21]: df.sort_index(axis=1, ascending=False)
Out[21]:
      D      C      B      A
2013-01-01 -1.135632 -1.509059 -0.282863 0.469112
2013-01-02 -1.044236 0.119209 -0.173215 1.212112
2013-01-03 1.071804 -0.494929 -2.104569 -0.861849
2013-01-04 0.271860 -1.039575 -0.706771 0.721555
2013-01-05 -1.087401 0.276232 0.567020 -0.424972
2013-01-06 0.524988 -1.478427 0.113648 -0.673690
```

Sorting by values:

```
In [22]: df.sort_values(by='B') Sort descending by 'D'
Out[22]:
      A      B      C      D
2013-01-03 -0.861849 -2.104569 -0.494929 1.071804
2013-01-04 0.721555 -0.706771 -1.039575 0.271860
2013-01-01 0.469112 -0.282863 -1.509059 -1.135632
2013-01-02 1.212112 -0.173215 0.119209 -1.044236
2013-01-06 -0.673690 0.113648 -1.478427 0.524988
2013-01-05 -0.424972 0.567020 0.276232 -1.087401
```

Selection

Note: While standard Python / Numpy expressions for selecting and setting are intuitive and come in handy for interactive work, for production code, we recommend the optimized pandas data access methods, `.at`, `.iat`, `.loc` and `.iloc`.

See the indexing documentation [Indexing and Selecting Data](#) and [MultiIndex / Advanced Indexing](#).

Getting

Selecting a single column, which yields a Series, equivalent to `df.A`:

```
In [23]: df['A']
Out[23]:
2013-01-01  0.469112
2013-01-02  1.212112
2013-01-03  -0.861849
2013-01-04   0.721555
2013-01-05  -0.424972
2013-01-06  -0.673690
Freq: D, Name: A, dtype: float64
```

Selecting via `[]`, which slices the rows.

```
In [24]: df[0:3] Select first 3 rows in column
```

```
Out[24]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804

```
In [25]: df['20130102':'20130104']
```

```
Out[25]:
```

	A	B	C	D
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860

Selection by Label

See more in [Selection by Label](#).

For getting a cross section using a label:

```
In [26]: df.loc[dates[0]]
```

```
Out[26]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
Name:	2013-01-01 00:00:00			
				dtype: float64

Selecting on a multi-axis by label:

```
In [27]: df.loc[:,['A','B']]
```

```
Out[27]:
```

	A	B
2013-01-01	0.469112	-0.282863
2013-01-02	1.212112	-0.173215
2013-01-03	-0.861849	-2.104569
2013-01-04	0.721555	-0.706771
2013-01-05	-0.424972	0.567020
2013-01-06	-0.673690	0.113648

Showing label slicing, both endpoints are *included*:

```
In [28]: df.loc['20130102':'20130104',['A','B']]
```

```
Out[28]:
```

	A	B
2013-01-02	1.212112	-0.173215
2013-01-03	-0.861849	-2.104569
2013-01-04	0.721555	-0.706771

Reduction in the dimensions of the returned object:

```
In [29]: df.loc['20130102',['A','B']]
```

```
Out[29]:
```

	A	B
2013-01-02	1.212112	-0.173215
Name:	-0.861849	-2.104569
	0.721555	-0.706771
		dtype: float64

For getting a scalar value:

```
In [30]: df.loc[dates[0],'A']
```

```
Out[30]: 0.46911229990718628
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [31]: df.at[dates[0],'A']
```

```
Out[31]: 0.46911229990718628
```

Selection by Position

See more in [Selection by Position](#).

Select via the position of the passed integers:

```
In [32]: df.iloc[3]
Out[32]:
A    0.721555
B   -0.706771
C   -1.039575
D    0.271860
Name: 2013-01-04 00:00:00, dtype: float64
```

By integer slices, acting similar to numpy/python:

```
In [33]: df.iloc[3:5,0:2]
Out[33]:
      A      B
2013-01-04  0.721555 -0.706771
2013-01-05 -0.424972  0.567020
```

By lists of integer position locations, similar to the numpy/python style:

```
In [34]: df.iloc[[1,2,4],[0,2]]
Out[34]:
      A      C
2013-01-02  1.212112  0.119209
2013-01-03 -0.861849 -0.494929
2013-01-05 -0.424972  0.276232
```

For slicing rows explicitly:

```
In [35]: df.iloc[1:3,:]
Out[35]:
      A      B      C      D
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804
```

For slicing columns explicitly:

```
In [36]: df.iloc[:,1:3]
Out[36]:
      B      C
2013-01-01 -0.282863 -1.509059
2013-01-02 -0.173215  0.119209
2013-01-03 -2.104569 -0.494929
2013-01-04 -0.706771 -1.039575
2013-01-05  0.567020  0.276232
2013-01-06  0.113648 -1.478427
```

For getting a value explicitly:

```
In [37]: df.iloc[1,1]
Out[37]: -0.17321464905330858
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [38]: df.iat[1,1]
Out[38]: -0.17321464905330858
```

Boolean Indexing

Using a single column's values to select data.

```
In [39]: df[df.A > 0]
Out[39]:
      A      B      C      D
2013-01-01  0.469112 -0.282863 -1.509059 -1.135632
2013-01-02  1.212112 -0.173215  0.119209 -1.044236
2013-01-04  0.721555 -0.706771 -1.039575  0.271860
```

Selecting values from a DataFrame where a boolean condition is met.

```
In [40]: df[df > 0]
Out[40]:
      A      B      C      D
2013-01-01  0.469112  NaN  NaN  NaN
2013-01-02  1.212112  NaN  0.119209  NaN
2013-01-03  NaN  NaN  NaN  1.071804
2013-01-04  0.721555  NaN  NaN  0.271860
2013-01-05  NaN  0.567020  0.276232  NaN
2013-01-06  NaN  0.113648  NaN  0.524988
```

Using the `isin()` method for filtering:

```
In [41]: df2 = df.copy()
In [42]: df2['E'] = ['one', 'one', 'two', 'three', 'four', 'three']
In [43]: df2
Out[43]:
      A      B      C      D      E
2013-01-01  0.469112 -0.282863 -1.509059 -1.135632  one
2013-01-02  1.212112 -0.173215  0.119209 -1.044236  one
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804  two
2013-01-04  0.721555 -0.706771 -1.039575  0.271860  three
2013-01-05 -0.424972  0.567020  0.276232 -1.087401  four
2013-01-06 -0.673690  0.113648 -1.478427  0.524988  three

In [44]: df2[df2['E'].isin(['two', 'four'])]
Out[44]:
      A      B      C      D      E
2013-01-03 -0.861849 -2.104569 -0.494929  1.071804  two
2013-01-05 -0.424972  0.567020  0.276232 -1.087401  four
```

Setting

Setting a new column automatically aligns the data by the indexes.

```
In [45]: s1 = pd.Series([1,2,3,4,5,6], index=pd.date_range('20130102', periods=6))
In [46]: s1
Out[46]:
2013-01-02    1
2013-01-03    2
2013-01-04    3
2013-01-05    4
2013-01-06    5
2013-01-07    6
Freq: D, dtype: int64

In [47]: df['F'] = s1
```

Setting values by label:

```
In [48]: df.at[dates[0], 'A'] = 0
```

Setting values by position:

```
In [49]: df.iat[0,1] = 0
```

Setting by assigning with a NumPy array:

```
In [50]: df.loc[:, 'D'] = np.array([5] * len(df))
```

The result of the prior setting operations.

```
In [51]: df
Out[51]:
      A      B      C      D      F
2013-01-01  0.000000  0.000000 -1.509059  5  NaN
2013-01-02  1.212112 -0.173215  0.119209  5  1.0
2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0
2013-01-05 -0.424972  0.567020  0.276232  5  4.0
2013-01-06 -0.673690  0.113648 -1.478427  5  5.0
```

A where operation with setting.

```
In [52]: df2 = df.copy()
In [53]: df2[df2 > 0] = -df2
In [54]: df2
Out[54]:
   A      B      C  D      F
2013-01-01  0.000000  0.000000 -1.509059 -5  NaN
2013-01-02 -1.212112 -0.173215 -0.119209 -5 -1.0
2013-01-03 -0.861849 -2.104569 -0.494929 -5 -2.0
2013-01-04 -0.721555 -0.706771 -1.039575 -5 -3.0
2013-01-05 -0.424972 -0.567020 -0.276232 -5 -4.0
2013-01-06 -0.673690 -0.113648 -1.478427 -5 -5.0
```

Missing Data

pandas primarily uses the value `np.nan` to represent missing data. It is by default not included in computations. See the [Missing Data section](#).

Reindexing allows you to change/add/delete the index on a specified axis. This returns a copy of the data.

```
In [55]: df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ['E'])
In [56]: df1.loc[dates[0]:dates[1],'E'] = 1
In [57]: df1
Out[57]:
   A      B      C  D      F      E
2013-01-01  0.000000  0.000000 -1.509059  5  NaN  1.0
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0  NaN
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0  NaN
```

To drop any rows that have missing data.

```
In [58]: df1.dropna(how='any')
Out[58]:
   A      B      C  D      F      E
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
```

Filling missing data.

```
In [59]: df1.fillna(value=5)  Fill value using random number
Out[59]: generation
   A      B      C  D      F      E
2013-01-01  0.000000  0.000000 -1.509059  5  5.0  1.0
2013-01-02  1.212112 -0.173215  0.119209  5  1.0  1.0
2013-01-03 -0.861849 -2.104569 -0.494929  5  2.0  5.0
2013-01-04  0.721555 -0.706771 -1.039575  5  3.0  5.0
```

To get the boolean mask where values are nan.

```
In [60]: pd.isna(df1)
Out[60]:
   A      B      C      D      F      E
2013-01-01  False  False  False  False  True  False
2013-01-02  False  False  False  False  False  False
2013-01-03  False  False  False  False  False  True
2013-01-04  False  False  False  False  False  True
```

Operations

See the [Basic section on Binary Ops](#).

Stats

Operations in general *exclude* missing data.

Performing a descriptive statistic:

```
In [61]: df.mean()
Out[61]:
A -0.004474
B -0.383981
C -0.687758
D 5.000000
F 3.000000
dtype: float64
```

Same operation on the other axis:

```
In [62]: df.mean(1)
Out[62]:
2013-01-01  0.872735
2013-01-02  1.431621
2013-01-03  0.707731
2013-01-04  1.395042
2013-01-05  1.883656
2013-01-06  1.592306
Freq: D, dtype: float64
```

Operating with objects that have different dimensionality and need alignment. In addition, pandas automatically broadcasts along the specified dimension.

```
In [63]: s = pd.Series([1,3,5,np.nan,6,8], index=dates).shift(2)

In [64]: s
Out[64]:
2013-01-01  NaN
2013-01-02  NaN
2013-01-03  1.0
2013-01-04  3.0
2013-01-05  5.0
2013-01-06  NaN
Freq: D, dtype: float64

In [65]: df.sub(s, axis='index')
Out[65]:
       A      B      C      D      F
2013-01-01  NaN    NaN    NaN    NaN    NaN
2013-01-02  NaN    NaN    NaN    NaN    NaN
2013-01-03 -1.861849 -3.104569 -1.494929  4.0  1.0
2013-01-04 -2.278445 -3.706771 -4.039575  2.0  0.0
2013-01-05 -5.424972 -4.432980 -4.723768  0.0 -1.0
2013-01-06  NaN    NaN    NaN    NaN    NaN
```

Apply

Applying functions to the data:

```
In [66]: df.apply(np.cumsum)
Out[66]:
       A      B      C      D      F
2013-01-01  0.000000  0.000000 -1.509059  5  NaN
2013-01-02  1.212112 -0.173215 -1.389850 10  1.0
2013-01-03  0.350263 -2.277784 -1.884779 15  3.0
2013-01-04  1.071818 -2.984555 -2.924354 20  6.0
2013-01-05  0.646846 -2.417535 -2.648122 25 10.0
2013-01-06 -0.026844 -2.303886 -4.126549 30 15.0

In [67]: df.apply(lambda x: x.max() - x.min())
Out[67]:
A  2.073961
B  2.671590
C  1.785291
D  0.000000
F  4.000000
dtype: float64
```

Histogramming

See more at [Histogramming and Discretization](#).

```
In [68]: s = pd.Series(np.random.randint(0, 7, size=10))
```

```
In [69]: s
```

```
Out[69]:
```

```
0    4  
1    2  
2    1  
3    2  
4    6  
5    4  
6    4  
7    6  
8    4  
9    4  
dtype: int64
```

```
In [70]: s.value_counts()
```

```
Out[70]:
```

```
4    5  
6    2  
2    2  
1    1  
dtype: int64
```

String Methods

Series is equipped with a set of string processing methods in the `str` attribute that make it easy to operate on each element of the array, as in the code snippet below. Note that pattern-matching in `str` generally uses [regular expressions](#) by default (and in some cases always uses them). See more at [Vectorized String Methods](#).

```
In [71]: s = pd.Series(['A', 'B', 'C', 'Aaba', 'Baca', np.nan, 'CABA', 'dog', 'cat'])
```

```
In [72]: s.str.lower()
```

```
Out[72]:
```

```
0    a  
1    b  
2    c  
3   aaba  
4   baca  
5    NaN  
6   cab  
7   dog  
8   cat  
dtype: object
```

Merge

Concat

pandas provides various facilities for easily combining together Series, DataFrame, and Panel objects with various kinds of set logic for the indexes and relational algebra functionality in the case of join / merge-type operations.

See the [Merging section](#).

Concatenating pandas objects together with `concat()`:

```
In [73]: df = pd.DataFrame(np.random.randn(10, 4))
```

```
In [74]: df
```

```
Out[74]:
```

```
   0   1   2   3  
0 -0.548702 1.467327 -1.015962 -0.483075  
1  1.637550 -1.217659 -0.291519 -1.745505  
2 -0.263952  0.991460 -0.919069  0.266046  
3 -0.709661  1.669052  1.037882 -1.705775  
4 -0.919854 -0.042379  1.247642 -0.009920  
5  0.290213  0.495767  0.362949  1.548106  
6 -1.131345 -0.089329  0.337863 -0.945867  
7 -0.932132  1.956030  0.017587 -0.016692  
8 -0.575247  0.254161 -1.143704  0.215897  
9  1.193555 -0.077118 -0.408530 -0.862495
```

```
# break it into pieces
```

```
In [75]: pieces = [df[3], df[3:7], df[7:]]
```

```
In [76]: pd.concat(pieces)
```

```
Out[76]:
```

```
   0   1   2   3  
0 -0.548702 1.467327 -1.015962 -0.483075  
1  1.637550 -1.217659 -0.291519 -1.745505  
2 -0.263952  0.991460 -0.919069  0.266046  
3 -0.709661  1.669052  1.037882 -1.705775  
4 -0.919854 -0.042379  1.247642 -0.009920  
5  0.290213  0.495767  0.362949  1.548106  
6 -1.131345 -0.089329  0.337863 -0.945867  
7 -0.932132  1.956030  0.017587 -0.016692  
8 -0.575247  0.254161 -1.143704  0.215897  
9  1.193555 -0.077118 -0.408530 -0.862495
```

Join

SQL style merges. See the [Database style joining](#) section.

```
In [77]: left = pd.DataFrame({'key': ['foo', 'foo'], 'lval': [1, 2]})
```

```
In [78]: right = pd.DataFrame({'key': ['foo', 'foo'], 'rval': [4, 5]})
```

```
In [79]: left
```

```
Out[79]:
```

```
   key  lval  
0 foo    1  
1 foo    2
```

```
In [80]: right
```

```
Out[80]:
```

```
   key  rval  
0 foo    4  
1 foo    5
```

```
In [81]: pd.merge(left, right, on='key')
```

```
Out[81]:
```

```
   key  lval  rval  
0 foo    1    4  
1 foo    1    5  
2 foo    2    4  
3 foo    2    5
```

Another example that can be given is:

```
In [82]: left = pd.DataFrame({'key': ['foo', 'bar'], 'lval': [1, 2]})
```

```
In [83]: right = pd.DataFrame({'key': ['foo', 'bar'], 'rval': [4, 5]})
```

```
In [84]: left
```

```
Out[84]:
```

```
   key  lval  
0  foo     1  
1  bar     2
```

```
In [85]: right
```

```
Out[85]:
```

```
   key  rval  
0  foo     4  
1  bar     5
```

```
In [86]: pd.merge(left, right, on='key')
```

```
Out[86]:
```

```
   key  lval  rval  
0  foo     1     4  
1  bar     2     5
```

Append

Append rows to a dataframe. See the [Appending](#) section.

```
In [87]: df = pd.DataFrame(np.random.randn(8, 4), columns=['A','B','C','D'])
```

```
In [88]: df
```

```
Out[88]:
```

```
       A      B      C      D  
0  1.346061  1.511763  1.627081 -0.990582  
1 -0.441652  1.211526  0.268520  0.024580  
2 -1.577585  0.396823 -0.105381 -0.532532  
3  1.453749  1.208843 -0.080952 -0.264610  
4 -0.727965 -0.589346  0.339969 -0.693205  
5 -0.339355  0.593616  0.884345  1.591431  
6  0.141809  0.220390  0.435589  0.192451  
7 -0.096701  0.803351  1.715071 -0.708758
```

```
In [89]: s = df.iloc[3]
```

```
In [90]: df.append(s, ignore_index=True)
```

```
Out[90]:
```

```
       A      B      C      D  
0  1.346061  1.511763  1.627081 -0.990582  
1 -0.441652  1.211526  0.268520  0.024580  
2 -1.577585  0.396823 -0.105381 -0.532532  
3  1.453749  1.208843 -0.080952 -0.264610  
4 -0.727965 -0.589346  0.339969 -0.693205  
5 -0.339355  0.593616  0.884345  1.591431  
6  0.141809  0.220390  0.435589  0.192451  
7 -0.096701  0.803351  1.715071 -0.708758  
8  1.453749  1.208843 -0.080952 -0.264610
```

Grouping

By “group by” we are referring to a process involving one or more of the following steps:

- **Splitting** the data into groups based on some criteria
- **Applying** a function to each group independently
- **Combining** the results into a data structure

See the [Grouping](#) section.

```
In [91]: df = pd.DataFrame({'A' : ['foo', 'bar', 'foo', 'bar',
...: 'foo', 'bar', 'foo', 'foo'],
...: 'B' : ['one', 'one', 'two', 'three',
...: 'two', 'two', 'one', 'three'],
...: 'C' : np.random.randn(8),
...: 'D' : np.random.randn(8))})
```

```
In [92]: df
```

```
Out[92]:
```

	A	B	C	D
0	foo	one	-1.202872	-0.055224
1	bar	one	-1.814470	2.395985
2	foo	two	1.018601	1.552825
3	bar	three	-0.595447	0.166599
4	foo	two	1.395433	0.047609
5	bar	two	-0.392670	-0.136473
6	foo	one	0.007207	-0.561757
7	foo	three	1.928123	-1.623033

Grouping and then applying the `sum()` function to the resulting groups.

```
In [93]: df.groupby('A').sum()
```

```
Out[93]:
```

	C	D
A		
bar	-2.802588	2.42611
foo	3.146492	-0.63958

Grouping by multiple columns forms a hierarchical index, and again we can apply the sum function.

```
In [94]: df.groupby(['A','B']).sum()  Group by A,B,C
```

```
Out[94]:
```

	C	D
A		
B		
bar one	-1.814470	2.395985
three	-0.595447	0.166599
two	-0.392670	-0.136473
foo one	-1.195665	-0.616981
three	1.928123	-1.623033
two	2.414034	1.600434

Reshaping

See the sections on [Hierarchical Indexing](#) and [Reshaping](#).

Stack

```
In [95]: tuples = list(zip(*[['bar', 'bar', 'baz', 'baz',
...: 'foo', 'foo', 'qux', 'qux'],
...: ['one', 'two', 'one', 'two',
...: 'one', 'two', 'one', 'two]]))
```

```
In [96]: index = pd.MultiIndex.from_tuples(tuples, names=['first', 'second'])
```

```
In [97]: df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=['A', 'B'])
```

```
In [98]: df2 = df[:4]
```

```
In [99]: df2
```

```
Out[99]:
```

	A	B
first	second	
bar	one	0.029399 -0.542108
	two	0.282696 -0.087302
baz	one	-1.575170 1.771208
	two	0.816482 1.100230

The `stack()` method “compresses” a level in the DataFrame’s columns.

```
In [100]: stacked = df2.stack()
```

```
In [101]: stacked
```

```
Out[101]:
```

```
first second
bar one A 0.029399
      B -0.542108
two  A 0.282696
      B -0.087302
baz  one A -1.575170
      B 1.771208
two  A 0.816482
      B 1.100230
dtype: float64
```

With a “stacked” DataFrame or Series (having a MultiIndex as the index), the inverse operation of `stack()` is `unstack()`, which by default unstacks the **last level**:

```
In [102]: stacked.unstack()
```

```
Out[102]:
```

```
A     B
first second
bar one 0.029399 -0.542108
      two 0.282696 -0.087302
baz  one -1.575170 1.771208
      two 0.816482 1.100230
```

```
In [103]: stacked.unstack(1)
```

```
Out[103]:
```

```
second   one    two
first
bar A 0.029399 0.282696
      B -0.542108 -0.087302
baz A -1.575170 0.816482
      B 1.771208 1.100230
```

```
In [104]: stacked.unstack(0)
```

```
Out[104]:
```

```
first   bar   baz
second
one  A 0.029399 -1.575170
      B -0.542108 1.771208
two  A 0.282696 0.816482
      B -0.087302 1.100230
```

Pivot Tables

See the section on [Pivot Tables](#).

```
In [105]: df = pd.DataFrame({'A' : ['one', 'one', 'two', 'three'] * 3,
```

```
...:           'B' : ['A', 'B', 'C'] * 4,
...:           'C' : ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 2,
...:           'D' : np.random.randn(12),
...:           'E' : np.random.randn(12)})
```

```
In [106]: df
```

```
Out[106]:
```

```
A  B  C  D  E
0  one A  foo 1.418757 -0.179666
1  one B  foo -1.879024 1.291836
2  two C  foo 0.536826 -0.009614
3  three A  bar 1.006160 0.392149
4  one B  bar -0.029716 0.264599
5  one C  bar -1.146178 -0.057409
6  two A  foo 0.100900 -1.425638
7  three B  foo -1.035018 1.024098
8  one C  foo 0.314665 -0.106062
9  one A  bar -0.773723 1.824375
10 two B  bar -1.170653 0.595974
11 three C  bar 0.648740 1.167115
```

We can produce pivot tables from this data very easily:

```
In [107]: pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
```

```
Out[107]:
```

```
C      bar    foo  
A   B  
one A -0.773723 1.418757  
     B -0.029716 -1.879024  
     C -1.146178 0.314665  
three A 1.006160    NaN  
       B    NaN -1.035018  
       C  0.648740    NaN  
two  A    NaN  0.100900  
     B -1.170653    NaN  
     C    NaN  0.536826
```

Time Series

pandas has simple, powerful, and efficient functionality for performing resampling operations during frequency conversion (e.g., converting secondly data into 5-minute data). This is extremely common in, but not limited to, financial applications. See the [Time Series section](#).

```
In [108]: rng = pd.date_range('1/1/2012', periods=100, freq='S')
```

```
In [109]: ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)
```

```
In [110]: ts.resample('5Min').sum()
```

```
Out[110]:
```

```
2012-01-01  25083  
Freq: 5T, dtype: int64
```

Time zone representation:

```
In [111]: rng = pd.date_range('3/6/2012 00:00', periods=5, freq='D')
```

```
In [112]: ts = pd.Series(np.random.randn(len(rng)), rng)
```

```
In [113]: ts
```

```
Out[113]:
```

```
2012-03-06  0.464000  
2012-03-07  0.227371  
2012-03-08 -0.496922  
2012-03-09  0.306389  
2012-03-10 -2.290613  
Freq: D, dtype: float64
```

```
In [114]: ts_utc = ts.tz_localize('UTC')
```

```
In [115]: ts_utc
```

```
Out[115]:
```

```
2012-03-06 00:00:00+00:00  0.464000  
2012-03-07 00:00:00+00:00  0.227371  
2012-03-08 00:00:00+00:00 -0.496922  
2012-03-09 00:00:00+00:00  0.306389  
2012-03-10 00:00:00+00:00 -2.290613  
Freq: D, dtype: float64
```

Converting to another time zone:

```
In [116]: ts_utc.tz_convert('US/Eastern')
```

```
Out[116]:
```

```
2012-03-05 19:00:00-05:00  0.464000  
2012-03-06 19:00:00-05:00  0.227371  
2012-03-07 19:00:00-05:00 -0.496922  
2012-03-08 19:00:00-05:00  0.306389  
2012-03-09 19:00:00-05:00 -2.290613  
Freq: D, dtype: float64
```

Converting between time span representations:

```
In [117]: rng = pd.date_range('1/1/2012', periods=5, freq='M')
```

```
In [118]: ts = pd.Series(np.random.randn(len(rng)), index=rng)
```

```
In [119]: ts
```

```
Out[119]:
```

```
2012-01-31 -1.134623  
2012-02-29 -1.561819  
2012-03-31 -0.260838  
2012-04-30 0.281957  
2012-05-31 1.523962  
Freq: M, dtype: float64
```

```
In [120]: ps = ts.to_period()
```

```
In [121]: ps
```

```
Out[121]:
```

```
2012-01 -1.134623  
2012-02 -1.561819  
2012-03 -0.260838  
2012-04 0.281957  
2012-05 1.523962  
Freq: M, dtype: float64
```

```
In [122]: ps.to_timestamp()
```

```
Out[122]:
```

```
2012-01-01 -1.134623  
2012-02-01 -1.561819  
2012-03-01 -0.260838  
2012-04-01 0.281957  
2012-05-01 1.523962  
Freq: MS, dtype: float64
```

Converting between period and timestamp enables some convenient arithmetic functions to be used. In the following example, we convert a quarterly frequency with year ending in November to 9am of the end of the month following the quarter end:

```
In [123]: prng = pd.period_range('1990Q1', '2000Q4', freq='Q-NOV')
```

```
In [124]: ts = pd.Series(np.random.randn(len(prng)), prng)
```

```
In [125]: ts.index = (prng.asfreq('M', 'e') + 1).asfreq('H', 's') + 9
```

```
In [126]: ts.head()
```

```
Out[126]:
```

```
1990-03-01 09:00 -0.902937  
1990-06-01 09:00 0.068159  
1990-09-01 09:00 -0.057873  
1990-12-01 09:00 -0.368204  
1991-03-01 09:00 -1.144073  
Freq: H, dtype: float64
```

Categoricals

pandas can include categorical data in a DataFrame. For full docs, see the [categorical introduction](#) and the [API documentation](#).

```
In [127]: df = pd.DataFrame({'id':[1,2,3,4,5,6], "raw_grade":['a', 'b', 'b', 'a', 'a', 'e']})
```

Convert the raw grades to a categorical data type.

```
In [128]: df["grade"] = df["raw_grade"].astype("category")
```

```
In [129]: df["grade"]
```

```
Out[129]:
```

```
0    a  
1    b  
2    b  
3    a  
4    a  
5    e  
Name: grade, dtype: category  
Categories (3, object): [a, b, e]
```

Rename the categories to more meaningful names (assigning to Series.cat.categories is inplace!).

```
In [130]: df["grade"].cat.categories = ["very good", "good", "very bad"]
```

Reorder the categories and simultaneously add the missing categories (methods under Series .cat return a new Series by default).

```
In [131]: df["grade"] = df["grade"].cat.set_categories(["very bad", "bad", "medium", "good", "very good"])
```

```
In [132]: df["grade"]
```

```
Out[132]:
```

```
0    very good  
1      good  
2      good  
3  very good  
4  very good  
5    very bad  
Name: grade, dtype: category  
Categories (5, object): [very bad, bad, medium, good, very good]
```

Sorting is per order in the categories, not lexical order.

```
In [133]: df.sort_values(by="grade")
```

```
Out[133]:
```

```
   id raw_grade    grade  
5   6       e  very bad  
1   2       b     good  
2   3       b     good  
0   1       a  very good  
3   4       a  very good  
4   5       a  very good
```

Grouping by a categorical column also shows empty categories.

```
In [134]: df.groupby("grade").size()
```

```
Out[134]:
```

```
grade  
very bad    1  
bad         0  
medium      0  
good        2  
very good   3  
dtype: int64
```

Plotting

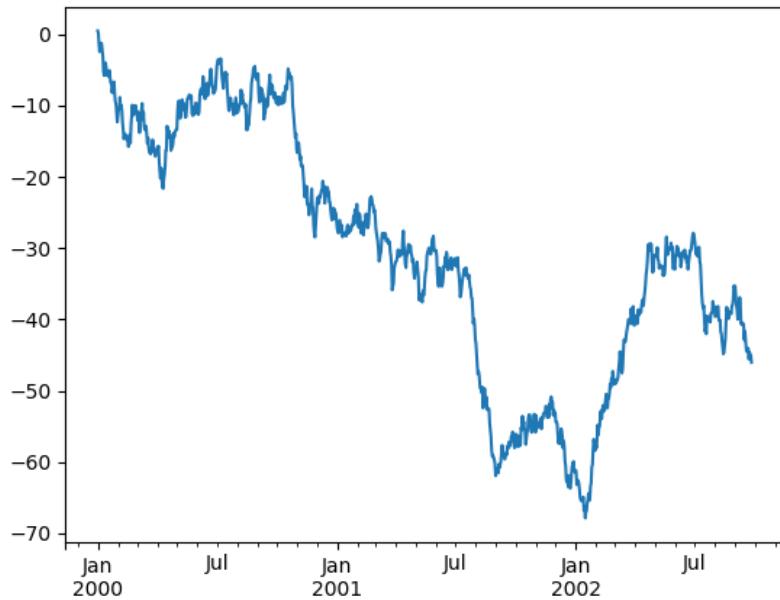
See the [Plotting](#) docs.

```
In [135]: ts = pd.Series(np.random.randn(1000), index=pd.date_range('1/1/2000', periods=1000))
```

```
In [136]: ts = ts.cumsum()
```

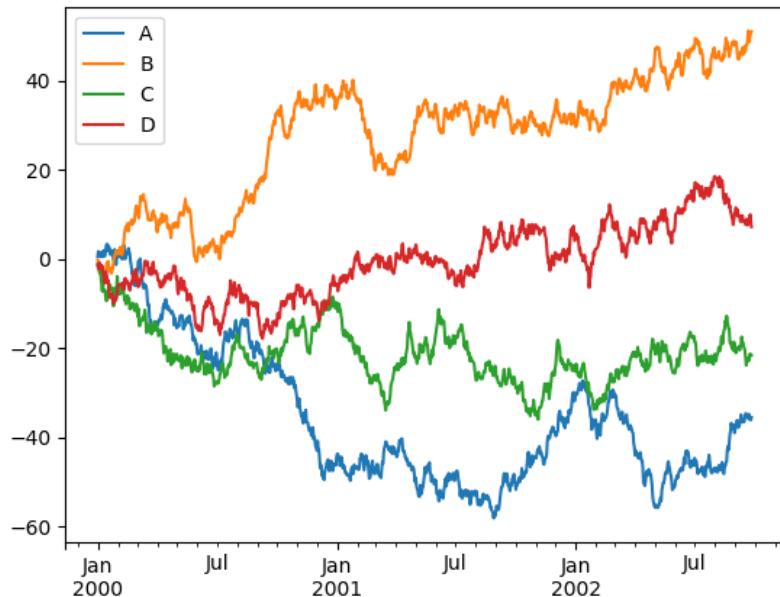
```
In [137]: ts.plot()
```

```
Out[137]: <matplotlib.axes._subplots.AxesSubplot at 0x7f213444c048>
```



On a DataFrame, the `plot()` method is a convenience to plot all of the columns with labels:

```
In [138]: df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index,
...:                      columns=['A', 'B', 'C', 'D'])
...:
In [139]: df = df.cumsum()
In [140]: plt.figure(); df.plot(); plt.legend(loc='best')
Out[140]: <matplotlib.legend.Legend at 0x7f212489a780>
```



Getting Data In/Out

CSV

Writing to a csv file.

```
In [141]: df.to_csv('foo.csv')
```

Reading from a csv file.

```
In [142]: pd.read_csv('foo.csv')
```

```
Out[142]:
```

```
UnNamed: 0      A      B      C      D
0  2000-01-01  0.266457 -0.399641 -0.219582  1.186860
1  2000-01-02 -1.170732 -0.345873  1.653061 -0.282953
2  2000-01-03 -1.734933  0.530468  2.060811 -0.515536
3  2000-01-04 -1.555121  1.452620  0.239859 -1.156896
4  2000-01-05  0.578117  0.511371  0.103552 -2.428202
5  2000-01-06  0.478344  0.449933 -0.741620 -1.962409
6  2000-01-07  1.235339 -0.091757 -1.543861 -1.084753
...
993 2002-09-20 -10.628548 -9.153563 -7.883146 28.313940
994 2002-09-21 -10.390377 -8.727491 -6.399645 30.914107
995 2002-09-22 -8.985362 -8.485624 -4.669462 31.367740
996 2002-09-23 -9.558560 -8.781216 -4.499815 30.518439
997 2002-09-24 -9.902058 -9.340490 -4.386639 30.105593
998 2002-09-25 -10.216020 -9.480682 -3.933802 29.758560
999 2002-09-26 -11.856774 -10.671012 -3.216025 29.369368
```

```
[1000 rows x 5 columns]
```

HDF5

Reading and writing to [HDFStores](#).

Writing to a HDF5 Store.

```
In [143]: df.to_hdf('foo.h5','df')
```

Reading from a HDF5 Store.

```
In [144]: pd.read_hdf('foo.h5','df')
```

```
Out[144]:
```

```
A      B      C      D
2000-01-01  0.266457 -0.399641 -0.219582  1.186860
2000-01-02 -1.170732 -0.345873  1.653061 -0.282953
2000-01-03 -1.734933  0.530468  2.060811 -0.515536
2000-01-04 -1.555121  1.452620  0.239859 -1.156896
2000-01-05  0.578117  0.511371  0.103552 -2.428202
2000-01-06  0.478344  0.449933 -0.741620 -1.962409
2000-01-07  1.235339 -0.091757 -1.543861 -1.084753
...
2002-09-20 -10.628548 -9.153563 -7.883146 28.313940
2002-09-21 -10.390377 -8.727491 -6.399645 30.914107
2002-09-22 -8.985362 -8.485624 -4.669462 31.367740
2002-09-23 -9.558560 -8.781216 -4.499815 30.518439
2002-09-24 -9.902058 -9.340490 -4.386639 30.105593
2002-09-25 -10.216020 -9.480682 -3.933802 29.758560
2002-09-26 -11.856774 -10.671012 -3.216025 29.369368
```

```
[1000 rows x 4 columns]
```

Excel

Reading and writing to [MS Excel](#).

Writing to an excel file.

```
In [145]: df.to_excel('foo.xlsx', sheet_name='Sheet1')
```

Reading from an excel file.

```
In [146]: pd.read_excel('foo.xlsx', 'Sheet1', index_col=None, na_values=['NA'])
```

```
Out[146]:
```

```
      A      B      C      D  
2000-01-01  0.266457 -0.399641 -0.219582  1.186860  
2000-01-02 -1.170732 -0.345873  1.653061 -0.282953  
2000-01-03 -1.734933  0.530468  2.060811 -0.515536  
2000-01-04 -1.555121  1.452620  0.239859 -1.156896  
2000-01-05  0.578117  0.511371  0.103552 -2.428202  
2000-01-06  0.478344  0.449933 -0.741620 -1.962409  
2000-01-07  1.235339 -0.091757 -1.543861 -1.084753  
...     ...     ...     ...  
2002-09-20 -10.628548 -9.153563 -7.883146  28.313940  
2002-09-21 -10.390377 -8.727491 -6.399645  30.914107  
2002-09-22 -8.985362 -8.485624 -4.669462  31.367740  
2002-09-23 -9.558560 -8.781216 -4.499815  30.518439  
2002-09-24 -9.902058 -9.340490 -4.386639  30.105593  
2002-09-25 -10.216020 -9.480682 -3.933802  29.758560  
2002-09-26 -11.856774 -10.671012 -3.216025  29.369368
```

```
[1000 rows x 4 columns]
```

Gotchas

If you are attempting to perform an operation you might see an exception like:

```
>>> if pd.Series([False, True, False]):  
    print("I was true")  
Traceback  
...  
ValueError: The truth value of an array is ambiguous. Use a.empty, a.any() or a.all().
```

See [Comparisons](#) for an explanation and what to do.

See [Gotchas](#) as well.